

Consensus Response to the Patched Master Collaborative Reference

Review prepared for Edward Wang’s Kitware internship record

July 15, 2026

Overall assessment

The patch is careful, constructive, and substantively correct. I accept the core scientific framing and almost all factual corrections. The remaining disagreements concern evidence labels and publication acceptance, not the paper’s thesis.

EEE detects reproducibility gaps. The follow-on work treats benchmark reproduction as layered system identification: it attempts to attribute and sometimes close disagreement by reconstructing recipe, deployment, execution-instrument, and artifact provenance, while recognizing that some historical measurements are non-identifiable from the surviving record examined.

The near-term manuscript should remain a hybrid position/systems-and-measurement paper. OLMo is the partially recoverable modern case; classic HELM supplies proxy-era and non-identifiable historical cases; a narrow Qwen study is optional rather than a dependency.

1 Decisions on the patch

A1. Commit 86ec84af

Accept the factual correction; modify the evidence label. The direct live repository observation is accepted, but the package does not include the Git object, a Git bundle, or an immutable remote reference. The master therefore uses *collaborator-verified*, *artifact not packaged*, rather than *Established*.

A2. Result stores

GPT-OSS: accept report freshness; reject preservation-only as the remaining step. A report timestamp does not establish run-artifact freshness, comparison freshness, or a complete provenance manifest. The cells are promising candidate results pending acceptance and preservation.

OLMo: accept. The reported local-halving signature makes the store stale relative to corrected comparison logic. Regenerate before citation.

RedPajama: accept with scope qualification. The live inspection strengthens the both-era proxy-execution claim. It does not prove either proxy is the original instrument. Copy, hash, and preserve the complete inputs and provenance.

Qwen: accept. No result store is present; the study remains pending and is not a near-term paper dependency.

A3. Probe scope and the 59% estimate

Accept the IFEval-only deployment-match scope, with a requirement to package sweep manifests and outputs. **Accept the 59% value only as an internal planning estimate.** Regenerate it from a frozen manifest with an explicit numerator, denominator, eligibility rule, and date before publication.

B1. Dense OLMo-2 HF divergence

Accept. Both genuine execution-instrument divergence and a dense-model probe artifact remain live explanations. The evidence locates a divergence, not its unique root cause.

B2. HuggingFaceClient wording

Accept with a source-version qualification. “Recorded as a HuggingFaceClient deployment” is precise; it must not imply a uniquely known historical package build.

B3. Drift ratio

Accept. The correct description is approximately 25–35x (about 30x), not two orders of magnitude.

B4. Canonical ledger

Accept and implement. The source of truth is `chronicle/data/master_claim_evidence_ledger_2026-07-15c.d`. Prose tables are snapshots. Store status is separately represented in `chronicle/data/store_status_ledger_2026-07-`

B5. Vendor-named front matter

Accept in principle. The front summary now refers to independent reconstructions. A concise evidence-status chapter remains near the front because it governs the chronology; detailed correspondence is preserved under `validation/`.

B6. Packaging and compilation

Accept and implement fully. The structured package contains both commit-ledger appendices, an accurate README, patch materials, canonical ledgers, and compiled primary documents.

2 Additional concerns from the independent review

Report freshness is not result freshness

Record separately: (1) run-artifact freshness, (2) comparison freshness, and (3) report freshness. A recent report timestamp establishes only the third unless a manifest links all three.

Add an orthogonal reproduction-target axis

The six-category taxonomy explains *why* executions disagree. The paper should also distinguish: artifact reconstruction, procedural reproduction, and claim-level robustness. Official released artifacts remain ground truth for the historical score; reruns test procedure stability and identifiability.

Scope non-identifiability

Use “non-identifiable from the surviving public and repository evidence examined,” not an absolute claim that no private archive could resolve the instrument.

Keep EvalAudit’s claim at the capability level

The current defensible claim is support for exact-spec replay, provenance preservation, layered comparison, and controlled diagnosis. Do not claim it guarantees reproducible evaluation or “fixes HELM” without system validation.

Preserve live-only evidence through attestations

For each live-store or post-archive Git claim, preserve identifiers and hashes, times, code and dirty state, run specs and targets, model/deployment revisions, environment and hardware identity, and generation/comparison commands.

3 Final consensus

There is no remaining substantive disagreement about the paper’s core thesis or near-term story.

A public LLM benchmark recipe is not a complete experimental record. Benchmark reproduction is a layered system-identification problem spanning recipe, deployment, execution instrument, and artifact interpretation. Controlled reconstruction can attribute and sometimes close apparent gaps, while some historical measurements remain non-identifiable from the surviving evidence examined.

The narrative is: EEE detects; EvalAudit replays, preserves, compares, and diagnoses; OLMo is partially recoverable pending ordinary-path confirmation; RedPajama compares declared proxy eras; GPT-J/GPT-NeoX/OPT illustrate origin-instrument non-identifiability; the paper derives a minimum provenance standard; and Qwen is an optional current-model demonstration.

The remaining work is operational: regenerate OLMo, preserve and audit GPT-OSS and RedPajama, package sweep manifests, complete or document failure of OLMo confirmation, regenerate denominators, and keep the canonical ledger synchronized.

With those qualifications incorporated, the chronology and paper strategy have reached consensus.